

CS553 Cloud Computing

Performance

Shruti Gupta (A20381966)

Sagar Mane (A20379756)

System Information:

The project aims to benchmarking different part of computer system, from the CPU, GPU, memory, disk, and network. This benchmarking implemented using KVM virtual machine m1.medium (2 virtual processors with 4GB RAM and 40GB disk) and CentOS 7 Linux for the OS.

```
[cc@pal-shruti-sagar ~]$ lscpu
Architecture:          x86_64
CPU op-mode(s):        32-bit, 64-bit
Byte Order:            Little Endian
CPU(s):                2
On-line CPU(s) list:   0,1
Thread(s) per core:    1
Core(s) per socket:    1
Socket(s):             2
NUMA node(s):         1
Vendor ID:             GenuineIntel
CPU family:            6
Model:                42
Model name:            Intel Xeon E312xx (Sandy Bridge)
Stepping:              1
CPU MHz:               2299.998
BogoMIPS:              4599.99
Hypervisor vendor:     KVM
Virtualization type:   full
L1d cache:             32K
L1i cache:             32K
L2 cache:              4096K
NUMA node0 CPU(s):    0,1
```

Benchmarking result and Analysis:

CPU Benchmarking:


Theoretical Peak Performance:

Number of cores * Average Frequency * cycle per seconds


= 2 * 2.2 * 8

=35.2 GFLOPS

CPU Integer Operation



CPU Integer Operation

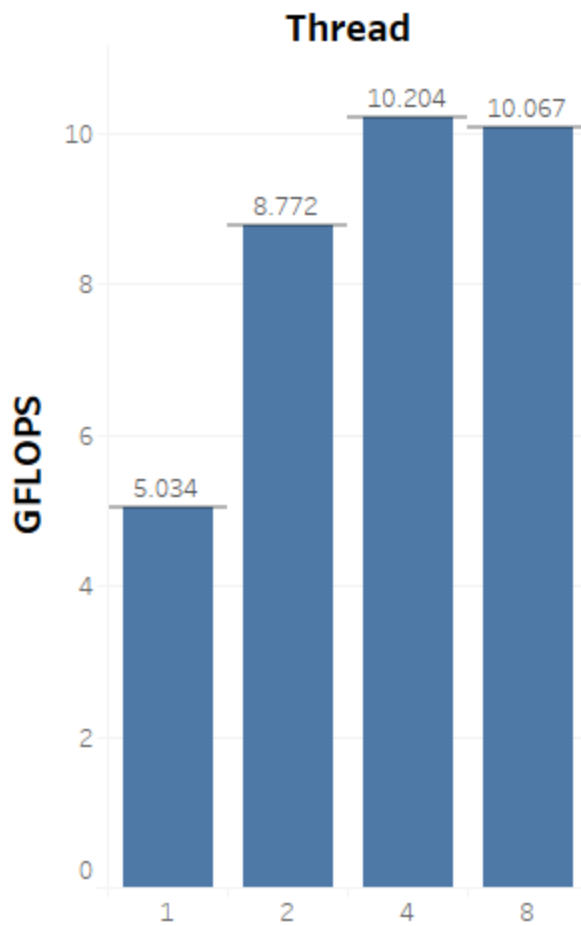


Above plot represents integer operation per seconds (GIOPS) for different threads and average and standard deviation of all threads.


We observed that as number of threads increase, we get more GIOPS. We get average 7.938 GIOPS and standard deviation 2.383 GIOPS.

Optimal number of thread align perfectly with the number of cores (2). Anything less, and the cores aren't being fully utilized. Anything more, and the threads are left waiting in queue to be swapped into context.

CPU Floating Point Operation



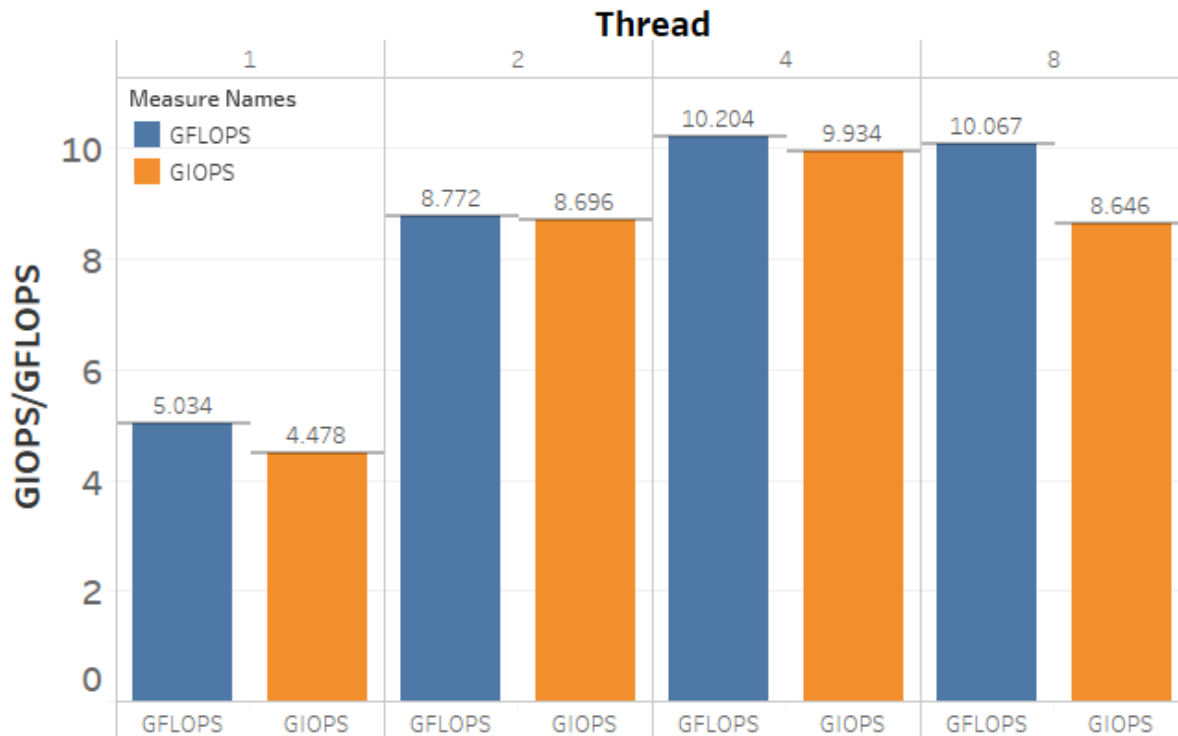
CPU Floating Point Operation



Above plot represents floating point operation per seconds (GFLOPS) for different threads and average and standard deviation of all threads.

We observed that as number of threads increase, we get more GFLOPS. We get average 8.519 GFLOPS and standard deviation 2.412 GFLOPS.

CPU Floating Point and Integer Operation



Above plot represents comparison between floating point operation and integer operation per seconds (GIOPS) for different threads.

For both we observed that GIOPS/GFLOPS increases with number of thread.

Efficiency with respect to theoretical peak performance:

$$=10.204/35.2$$

$$=29.15\%$$

For AVX instruction we can get more efficiency (76.42%) as compare to CPU operations.

LINPACK output:

```
[cc@pal-shruti-sagar linpack]$ ./xlinpack_xeon64
Input data or print help ? Type [data]/help :
```

```
Number of equations to solve (problem size): 5000
Leading dimension of array: 10000
Number of trials to run: 10
Data alignment value (in Kbytes): 400000
Current date/time: Thu Oct  5 03:30:44 2017
```

```
CPU frequency:      3.095 GHz
Number of CPUs: 2
Number of cores: 2
Number of threads: 2
```

Parameters are set to:

```
Number of tests                               : 1
Number of equations to solve (problem size)  : 5000
Leading dimension of array                    : 10000
Number of trials to run                      : 10
Data alignment value (in Kbytes)             : 400000
```

Maximum memory requested that can be used = 809800000, at the size = 5000

===== Timing linear equation system solver =====

Size	LDA	Align.	Time(s)	GFlops	Residual	Residual(norm)
5000	10000	400000	2.225	37.4831	2.581643e-11	3.599893e-02
5000	10000	400000	2.232	37.3650	2.581643e-11	3.599893e-02
5000	10000	400000	2.244	37.1524	2.581643e-11	3.599893e-02
5000	10000	400000	2.242	37.1988	2.581643e-11	3.599893e-02
5000	10000	400000	2.217	37.6024	2.581643e-11	3.599893e-02
5000	10000	400000	2.156	38.6745	2.581643e-11	3.599893e-02
5000	10000	400000	2.183	38.1978	2.581643e-11	3.599893e-02
5000	10000	400000	2.206	37.7974	2.581643e-11	3.599893e-02
5000	10000	400000	2.203	37.8510	2.581643e-11	3.599893e-02
5000	10000	400000	2.195	37.9946	2.581643e-11	3.599893e-02

Performance Summary (GFlops)

Size	LDA	Align.	Average	Maximal
5000	10000	400000	37.7317	38.6745

End of tests

Using LINPACK benchmarking we get 37.7317 GFLOPS.


Efficiency with respect to LINPACK benchmark:

=8.772/37.7317 =23.24%

This is not best performance as compare to LINPACK benchmark. We get more better performance if we performing operation using AVX instruction (did VAX implementation for extra credit).

Disk Benchmarking

Write Sequential



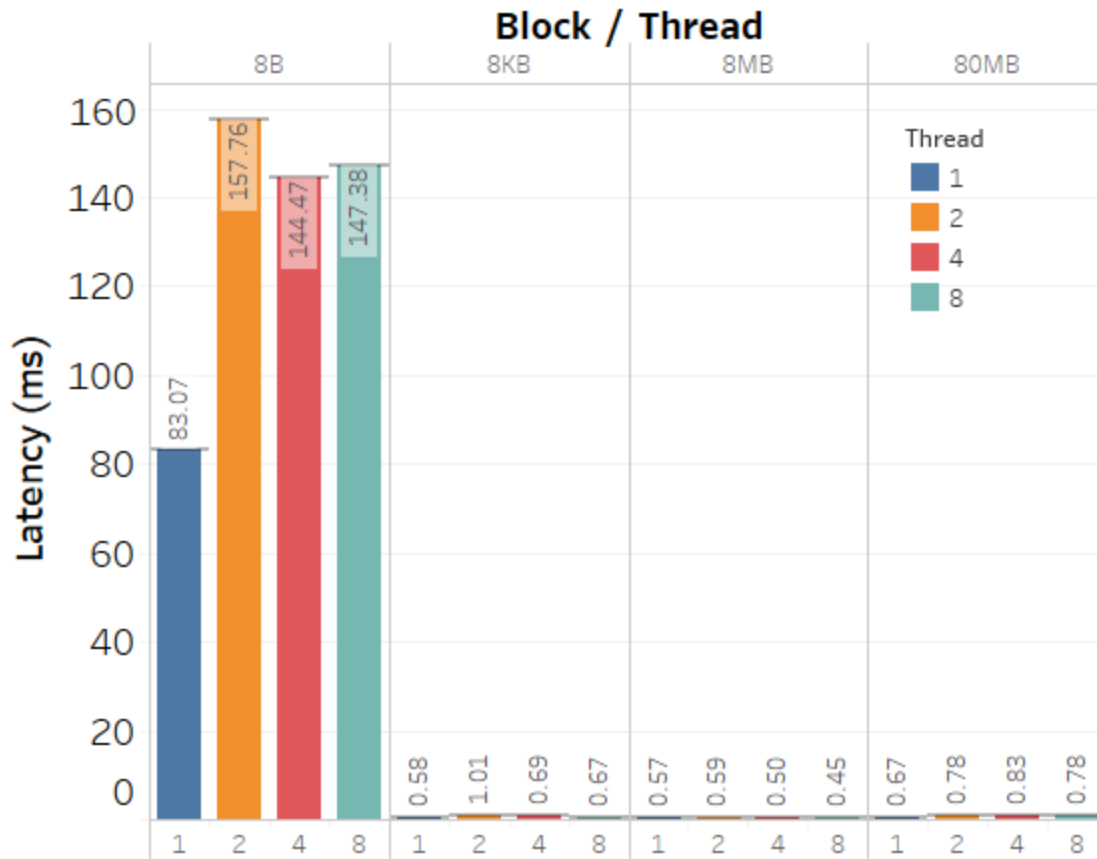
Write Sequential

Thread	≡	Average	Standard Deviation
1		1,220.0	814.9
2		973.4	705.4
4		1,138.1	822.9
8		1,220.0	901.1

Above plot represents the throughput for sequential writing on disk for different block size and number of threads and average and standard deviation of all threads. Throughput is measure in the MegaBytes per seconds.

We observed that we get low throughput for block size 8B as it takes more time to write file compare to other block size.


Write Sequential



Above plot represents the latency for sequential writing on disk for different block size and number of threads. latency is measure in milliseconds.

We observed that we get high latency for block size 8B as it takes more time to write file compare to other block size.

Write Random




Write Random

Thread	≡	Average	Standard Deviation
1		1,272.7	871.6
2		927.3	689.2
4		1,013.6	726.2
8		1,132.4	807.8

Above plot represents the throughput for random writing on disk for different block size and number of threads and average and standard deviation of all threads. Throughput is measure in the MegaBytes per seconds.

We observed that we get low throughput for block size 8B as it takes more time to write file compare to other block size.


Write Random



Above plot represents the latency for random writing on disk for different block size and number of threads. latency is measure in milliseconds.

We observed that we get high latency for block size 8B as it takes more time to write file compare to other block size.

Read Sequential




Read Sequential

Thread	≡	Average	Standard Deviation
1		43,756	42,688
2		15,629	13,760
4		9,221	6,923
8		4,987	3,630

Above plot represents the throughput for sequential reading on disk for different block size and number of threads and average and standard deviation of all threads. Throughput is measure in the MegaBytes per seconds.

We observed that we get low throughput for block size 8B as it takes more time to read file compare to other block size.


Read Sequential



Above plot represents the latency for sequential reading on disk for different block size and number of threads. latency is measure in milliseconds.

We observed that we get high latency for block size 8B as it takes more time to read file compare to other block size.

Read Random




Read Random

Thread	μ	Average	Standard Deviation
1		40,628	44,921
2		22,759	23,083
4		7,163	5,844
8		4,382	3,993

Above plot represents the throughput for random reading on disk for different block size and number of threads and average and standard deviation of all threads. Throughput is measure in the MegaBytes per seconds.

We observed that we get low throughput for block size 8B as it takes more time to read file compare to other block size.


Read Random



Above plot represents the latency for random reading on disk for different block size and number of threads. latency is measure in milliseconds.

We observed that we get high latency for block size 8B as it takes more time to read file compare to other block size.

Read + Write




Read + Write

Thread	≡	Average	Standard Deviation
1		1,199.2	795.1
2		874.4	604.6
4		1,025.8	695.5
8		947.4	650.8

Above plot represents the throughput for reading and writing file on disk for different block size and number of threads and average and standard deviation of all threads. Throughput is measure in the MegaBytes per seconds.

We observed that we get low throughput for block size 8B as it takes more time to read and write file compare to other block size. For more concurrency 8 we get better throughput.

Read + Write



Above plot represents the latency for reading and writing on disk for different block size and number of threads. latency is measure in milliseconds.

We observed that we get high latency for block size 8B as it takes more time to read and write file compare to other block size.

We observed that for concurrency 8 we get optimal number of performance as compare to another concurrency.


```
[cc@pal-shruti-sagar current]$ ./iozone -a
Iozone: Performance Test of File I/O
Version $Revision: 3.471 $
Compiled for 64 bit mode.
Build: linux
```

Run began: Thu Oct 5 03:53:09 2017

```
Auto Mode
Command line used: ./iozone -a
Output is in kBytes/sec
Time Resolution = 0.000001 seconds.
Processor cache size set to 1024 kBytes.
Processor cache line size set to 32 bytes.
File stride size set to 17 * record size.
```

	kB	reclen	write	rewrite	read	reread	random	random	bkwd	record	stride	write	rewrite	freed	freeread
64	4	695778	1421755	5860387	4564786	4274062	1991276	2772930	1991276	3057153	1183548	1690338	2772930	4274062	
64	8	1143223	2203800	5860307	6271021	5860307	1780008	2561267	2379626	3791156	2133730	1484662	3541098	2801873	
64	16	1049372	2067979	7100397	6421025	2467108	2923952	2467108	2561267	2006158	2772930	4274062	7100397		
64	32	1066042	2561267	7100397	9006179	7100397	2662899	3791156	1828508	3738358	2298136	2892445	4643754	7100397	
64	64	1017549	2467108	9006179	9318832	7100397	1734015	2923952	2561267	3791156	2467108	2923952	4274062	5735102	
128	4	743788	1732927	3560167	5784891	4267461	2035099	2843510	2660335	3445772	1967960	1684006	3657016	4407601	
128	8	1142750	2409592	6727225	7176872	5784891	2511022	4012317	3199360	5603747	2367096	1939522	4889281	6114306	
128	16	1152564	2453642	6727225	7082197	6114306	2843510	4135958	3106789	4717434	2464907	2166499	4934216	6114306	
128	32	1281909	2608629	6466138	7582312	6727225	3036502	4012317	3445772	4934216	2511022	3124872	5325799	6406138	
128	64	1294270	2714132	7582312	7917784	7082197	3199360	4135958	3468038	4934216	3236073	4135958	4889281	6727225	
128	128	1391557	2453642	3895854	8548124	6727225	3606777	4267461	3380677	4889281	3277486	3395523	5603747	7082197	
256	4	1080434	1911233	5117791	5022044	3197509	1907837	3316006	2692393	3197509	2097948	2131261	4422226	3950402	
256	8	1332507	2187712	6936061	6554972	5569035	2765777	4332998	3879045	5117791	2533571	2205688	5687020	5320671	
256	16	1256123	2097948	5810912	5841722	5217259	2148318	4332998	3326279	5320671	2148318	2533571	5117791	5428265	
256	32	1241598	2563818	6213578	6398720	5810112	2795955	3511189	3499745	5428265	2588541	2639446	5117791	6073004	
256	64	1274008	2872459	7120034	7364196	6719046	3368013	5217259	3935921	6107584	2808164	3368013	5569935	6891544	
256	128	1471220	3009318	6891544	6107548	5687020	3454704	5422734	3285566	5117791	2463808	3756894	5687020	6595234	
256	256	1235882	2783115	6398720	6554972	6398720	2880164	4907284	2692393	5569035	3285556	3368013	5428265	6249745	
512	4	1002183	1619757	4228947	4340054	3905895	1619757	2764505	2391924	3659616	1859721	1910902	4038998	3970896	
512	8	1125040	2265742	6736026	5822804	5624545	2668325	5019764	3304808	4457157	2600470	2560168	4973263	6472112	
512	16	1319250	2753870	6571132	6736026	4701087	3046915	5227492	3822466	5453155	2497638	2509312	5227492	6414119	
512	32	1004057	2845081	6571132	6736026	6319739	3086326	5019764	3504346	6319739	2497638	2521095	5278892	6551087	
512	64	1551867	2584820	6821616	6931711	7022379	3029721	5886650	3220553	6736026	2546058	3158966	4973263	6821616	
512	128	1546280	3012719	6736026	6821616	6472112	3653390	4882800	4271081	5227492	3325278	3510074	5566231	6714737	
512	256	1365383	2246778	5495016	5566231	5384786	2625909	5164632	3136320	5509113	2124526	4131319	5008059	5886650	
512	512	1433753	2722449	6018636	6018636	5822804	2922519	5214798	3029721	5439343	3201349	2722449	4650188	5951911	

1024	4	982873	1615374	4146499	4270175	3425544	1777861	3379719	2301937	3229770	1868347	1723636	3866533	3567824
1024	8	1230718	2169388	6128613	6393168	5336651	2572135	4995276	3951918	5120336	2270308	21881177	4943530	4720752
1024	16	1253344	2410457	5845038	6479797	4826859	2926114	4594582	3544270	4972145	2565989	2480051	5506686	6244448
1024	32	1212991	2409105	5250041	5529011	5226256	3020783	4942226	3955557	5906934	2730767	2730767	4926518	6244448
1024	64	1395018	2565989	6559147	6650556	5853803	3379719	5144870	3908759	6569179	2835325	2735327	5587740	9550309
1024	128	1406431	2798378	5690162	4995276	6128613	3251778	4943530	3966517	5950309	2746483	2993413	5821271	5623111
1024	256	1298041	2290886	4720752	5144870	5917516	3037876	5096034	3518141	5251818	3210456	3239514	5169641	5798981
1024	512	1389594	2455943	5564289	5630486	6025439	3083680	5506866	3433760	5782087	2645000	3891053	5065980	6025439
1024	1024	1416169	2723840	6208343	5813392	4317393	3314514	4995276	3094790	5728107	2852271	2900425	5251818	6016997
2048	4	958291	1959469	4741090	4686376	3618634	1972971	2964289	2412357	3319389	1882186	1785550	4257067	4172290
2048	8	1226291	2356760	6185123	5447681	4945823	2537792	4656281	3579429	3946119	2364455	2343259	5868220	5389574
2048	16	1257897	2547756	5719737	5900467	5349298	2901215	5674397	3611028	5430461	2429413	2458618	5758524	5290001
2048	32	1344739	2990975	6711905	5549749	5689430	3074635	5800851	3644736	5029807	2590122	2642401	6800553	5319484
2048	64	1454376	2916978	5970183	6167360	5280246	3121562	5521212	4385291	6176360	2805512	2656292	6075753	6092991
2048	128	1423398	2723672	5640860	5332694	5754221	3448675	5578583	3929871	6583305	2691246	2771120	10339282	11917129
2048	256	2528082	6041567	9567698	12561952	8608824	7293167	5119743	3915540	5785224	2809182	2855880	5489457	5754221
2048	512	1450315	2756890	5674397	6024617	5884299	3225896	5820505	3443127	10780463	5416763	6280083	6114678	9708272
1024	1024	1729816	5461535	11502234	11770166	12118884	6397004	9441504	6629029	11015480	4697018	6780783	11579763	12561952
2048	2048	1867050	6253312	12979541	10490890	11564174	6623917	4697018	3341339	5986827	3340044	3210223	5820505	5800851
4096	4	1019951	2032926	4946971	5050220	4100397	2160499	3949571	2834585	3957760	1813565	2140045	4737369	6433432
4096	8	1405213	2588965	6041154	5468151	4727091	2353791	4681411	3540193	4964124	2482709	2422154	5454263	5841899
4096	16	1329011	2068620	6291135	5497899	5424983	2923334	5589120	3660115	5580043	2587405	2556602	6041154	6252028
4096	32	1259539	2551665	5971855	6178018	5291314	3000952	4887857	3761077	5404504	2612587	2622558	6060333	5657127
4096	64	1396532	3119752	6522830	6545196	6693092	3747131	6142675	4638436	6311938	2996241	3002525	6480999	6471234
4096	128	1538614	3214912	6389405	6142675	6420447	3775126	6015769	4718699	5794602	2917377	2925823	5681538	6522830
4096	256	1520906	2946902	6401308	6468797	6449370	3673419	6060333	4043458	6628532	2927818	2866556	5736555	6439700
4096	512	1444681	3040789	6263611	6216019	6234064	3625358	6042207	4132949	6272759	3045101	3066750	5810280	6169296
4096	1024	1474944	3080040	6815229	6815229	6769576	3853020	6439700	4432592	6439700	3424463	3706707	5759634	6782940
4096	2048	1599346	3311586	6727163	7038605	6772244	3882626	5961494	4266367	6716643	2561939	4192455	5596402	6225028
4096	4096	1447358	6015769	11490987	11907264	10612681	6615769	11252164	7038605	12606249	6487827	7158858	10860922	12450933

32768	64	1515634	2298356	4949757	5413581	5985025	2986257	5017337	4448034	5556724	2373983	2203048	6081691	6377439
32768	128	1373803	2428129	4431558	5219305	5498840	3074847	4979166	4554909	5435849	2536625	2191107	5271585	5586993
32768	256	1569425	2316563	4743363	4967827	5578376	3237910	5300423	4808280	5885077	2777397	2123533	4512882	4875484
32768	512	1493190	2250669	4525364	4995909	5270544	3035191	4741431	4413184	5670203	2687251	2154222	4584231	4989770
32768	1024	1429231	2301781	4617033	5064669	5641806	3204690	4784345	4078572	5639722	2544279	2146048	4560047	4938020
32768	2048	1483183	2256915	4625579	4976822	5424599	2855351	4714918	4075186	6354146	2442282	2101840	4796199	5254625
32768	4096	1547477	2263270	4766583	5115186	5552908	2965190	5140055	4660400	6835835	2608503	2131107	4507797	4943881
32768	8192	1425067	2273491	4892841	5360374	5382416	3002633	4681832	3985377	6465644	2871338	2471532	4542264	4837218
32768	16384	1395520	2276427	6422439	6553510	6534193	3688863	6057032	4583467	5586085	3783819	3790812	4593270	6289304
65536	64	2136122	2315194	5366543	5609546	5913157	2850874	4592066	4518208	5514121	2525388	2224510	5129677	5413987
65536	128	1630289	2361234	4938623	5249187	5392532	2983917	4903472	5202492	5894644	2709897	2198604	5108038	5564802
65536	256	1744560	2326047	4642987	4798694	5130060	2940285	4615545	4752728	5432390	2586993	2157990	4546381	4743461
65536	512	1627085	2340147	4657463	4981855	5279635	3035888	4600251	4511312	5303574	2543615	2157990	4644321	5028425
65536	1024	1706068	2374841	4698221	4919589	5232325	2982557	4674809	4770379	5682958	2644493	2217656	4747885	4975003
65536	2048	1700684	2457123	4720975	4985560	4927544	2611497	4747229	4409348	4733495	2201862	2108690	4801880	5114787
65536	4096	1708443	2178767	4626519	4960279	5126042	2819928	4852658	4954022	7856167	2676674	2098099	4766988	5011556
65536	8192	1647338	2548544	5068573	5581866	5789486	3138527	5001981	5115993	6975037	2550436	2269021	4830571	5010004
65536	16384	1550601	2323648	4641231	4858111	4442485	2608820	4207086	3285662	4436820	2736252	2479799	4316427	4768972

131072	64	1769998	2491474	5159322	5853210	6188796	3063785	5071650	5992946	6448460	3058007	2376173	5231793	5891410
131072	128	1958107	3638169	7100696	7401446	7416523	3773157	4970496	6151884	6227001	2959609	2454488	5003520	5517653
131072	256	1871389	2509122	5040819	5694366	5905523	3371973	4912534	5821538	5967057	2942186	2377334	5125985	5457456
131072	512	1957584	2435006	4359895	4817441	4984105	3040620	4413801	5646234	5298146	2803206	2303435	4417525	4730636
131072	1024	1919626	2415311	4419301	4900797	5195501	3200854	5070855	5539112	6178294	2939889	2431215	5050729	5492680
131072	2048	2076549	2523519	5059142	5614519	6007550	3284265	5083891	5887687	6368235	2819060	2283847	5306328	5865826
131072	4096	2119256	2548400	5219920	4826280	4795379	2972089	4862092	5591847	6303396	2709213	2164582	4914202	5359298
131072	8192	1989225	2554463	5008899	5553660	6068962	3488280	5245370	5963821	8813465	2820637	2277611	5152601	5613888
131072	16384	1977099	2374541	4699096	5217196	5387026	3526182	4491256	3898955	5681243	2898977	2312620	4643490	5075958
262144	64	2148424	2827599	6223610	6912350	7509771	3326824	5580036	8638743	8029863	3492976	2779451	5773599	6651017
262144	128	2419732	2860772	5671070	6523092	7121890	3238870	6169551	9974269	7779207	3434076	2706391	5974249	6602134
262144	256	2362083	2641115	6126716	7540880	7513209	4020860	5281547	7966447	7041795	3312951	2723884	5280786	6028361
262144	512	2264234	2891383	5383513	6237946	7246558	3906363	5429424	7343795	7559804	3392310	2600893	7656936	7733519
262144	1024	2372107	2758947	6068453	5209948	6177246	3700716	5689263	8029629	7400988	3223374	2761857	5700028	6375576
262144	2048	2394402	2950411	7763278	7710578	7392479	3730726	5658003	8288829	7438188	2967940	2844534	6155769	7468249
262144	4096	1972222	2937319	5640291	6467418	7119399	3720035	5805548	7897326	7413163	2833436	2410079	5455663	6236177
262144	8192	2211456	3108073	5449417	6553418	5376537	3395872	5660384	7590229	7682992	2909154	2391257	5666014	6366900
262144	16384	2160285	3127123	5706862	5714426	4513264	2952867	5131123	4475967	4595798	2577366	2365259	4795208	5394686
524288	64	2101033	3322907	8042221	7832540	7666701	3310755	5750286	10330803	7841114	3359098	3055680	6258661	7546775
524288	128	1824511	2614564	6445009	6190755	7460086	3686798	5759126	10188644	7590015	3249763	3076319	6348864	7869370
524288	256	1984202	2691764	5178256	6627580	5184042	3209813	5333114	8185707	7127933	3155056	2951923	5920953	6011647
524288	512	2020853	2828606	5632847	6295716	4809985	2976980	5568644	8713737	6467890	3183775	2715001	5376429	6216885
524288	1024	2276353	2982451	5279681	6254638	6376183	3110152	5605323	8211505	7130938	3112295	2921995	5434721	6504488
524288	2048	1943888	2793670	7106006	7214647	7155581	3377403	5866020	7913548	7311632	2872592	2572321	6208040	7355630
524288	4096	1941403	3008915	6181758	7100705	7103802	3657578	5949757	7985854	6093185	2459898	2611319	6091834	7431447
524288	8192	2093050	2819814	6229584	7431748	7649952	3666066	5928232	8121585	5526450	2654219	2611620	6228401	6655260
524288	16384	2195427	2759483	5395306	5737323	5540290	3101672	5065685	4806494	5437261	2582674	2469864	5306922	5851939

iozone test complete.

IOZone benchmark for 1MB write sequential is 2071885Kb/s and we get 1569 Mb/sec

Efficiency is

= 1667/2071.885

=80.45%

IOZone benchmark for 1MB write random is 4031515 KB/s and we get 1616 Mb/sec

Efficiency is

= 1667/2071.885


=40.08%

We observed that for sequential write we get better performance as compare to random write.

According to result disk we are evaluating is spinning hard drive (HDD).

Memory Benchmarking:

Write Sequential




Write Sequential

Thread	Average	Standard Deviation
1	3,335	1,274
2	5,831	2,157
4	6,364	2,385
8	7,059	3,169

Above plot represents the throughput for sequential writing on memory for different block size and number of threads and average and standard deviation of all threads. Throughput is measure in the MegaBytes per seconds.

We observed that we get low throughput for block size 8B as it takes more time to write memory as compare to other block size.


Write Sequential



Above plot represents the latency for sequential writing on memory for different block size and number of threads. latency is measure in microseconds (us).

We observed that we get high latency for block size 8B as it takes more time to write memory compare to other block size.

Write Random



Write Random


Thread	Average	Standard Deviation
1	3,232	2,339
2	5,295	3,749
4	5,645	4,036
8	6,269	4,768

Above plot represents the throughput for random writing on memory for different block size and number of threads threads and average and standard deviation of all threads. Throughput is measure in the MegaBytes per seconds.

We observed that we get low throughput for block size 8B as it takes more time to write memory as compare to other block size.

We also observed that random throughput of random write with 8B is low as compare to sequential write.


Write Random



Above plot represents the latency for random writing on memory for different block size and number of threads. latency is measure in microseconds (us).

We observed that we get high latency for block size 8B as it takes more time to write memory compare to other block size.

Read Sequential




Read Sequential

Thread	Average	Standard Deviation
1	11,512	8,769
2	13,187	9,302
4	15,876	13,068
8	15,216	13,322

Above plot represents the throughput for sequential reading on memory for different block size and number of threads threads and average and standard deviation of all threads. Throughput is measure in the MegaBytes per seconds.

We observed that we get low throughput for block size 8B as it takes more time to read memory as compare to other block size.


Read Sequential



Above plot represents the latency for sequential reading on memory for different block size and number of threads. latency is measure in microseconds (us).

We observed that we get high latency for block size 8B as it takes more time to read memory compare to other block size.

Read Random



Read Random


Thread	Average	Standard Deviation
1	7,897	6,467
2	10,601	7,555
4	11,638	8,308
8	10,770	8,434

Above plot represents the throughput for random reading on memory for different block size and number of threads threads and average and standard deviation of all threads. Throughput is measure in the MegaBytes per seconds.

We observed that we get low throughput for block size 8B as it takes more time to read memory as compare to other block size.

We also observed that random throughput of random read with 8B is low as compare to sequential read.


Read Random



Above plot represents the latency for random reading on memory for different block size and number of threads. latency is measure in microseconds (us).

We observed that we get high latency for block size 8B as it takes more time to read memory compare to other block size.

Read + Write




Read + Write

Thread	Average	Standard Deviation
1	2,608	1,142
2	4,248	1,764
4	4,370	1,844
8	4,482	1,900

Above plot represents the throughput for reading and writing on memory for different block size and number of threads threads and average and standard deviation of all threads. Throughput is measure in the MegaBytes per seconds.

We observed that we get low throughput for block size 8B as it takes more time to read and write memory as compare to other block size.

Read + Write



Above plot represents the latency for reading and writing on memory for different block size and number of threads. latency is measure in microseconds (us).

We observed that we get high latency for block size 8B as it takes more time to read and write memory compare to other block size.

We observed that for concurrency 8 we get optimal number of performance as compare to another concurrency.

Stream Benchmark Output:

```
-----
STREAM version $Revision: 5.10 $
-----
This system uses 8 bytes per array element.
-----
Array size = 10000000 (elements), Offset = 0 (elements)
Memory per array = 76.3 MiB (= 0.1 GiB).
Total memory required = 228.9 MiB (= 0.2 GiB).
Each kernel will be executed 10 times.
  The *best* time for each kernel (excluding the first iteration)
  will be used to compute the reported bandwidth.
-----
Your clock granularity/precision appears to be 1 microseconds.
Each test below will take on the order of 28862 microseconds.
  (= 28862 clock ticks)
Increase the size of the arrays if this shows that
you are not getting at least 20 clock ticks per test.
-----
WARNING -- The above is only a rough guideline.
For best results, please be sure you know the
precision of your system timer.
-----
Function      Best Rate MB/s  Avg time     Min time     Max time
Copy:         4949.4    0.034250    0.032327    0.037157
Scale:        4827.4    0.035421    0.033144    0.039901
Add:          7147.6    0.036772    0.033578    0.042070
Triad:        6652.5    0.040410    0.036077    0.044236
-----
Solution Validates: avg error less than 1.000000e-13 on all three arrays
-----
```

Stream Benchmark for memory is 4949.4 MB/s and we get average 3335 Mb/s for sequential write.

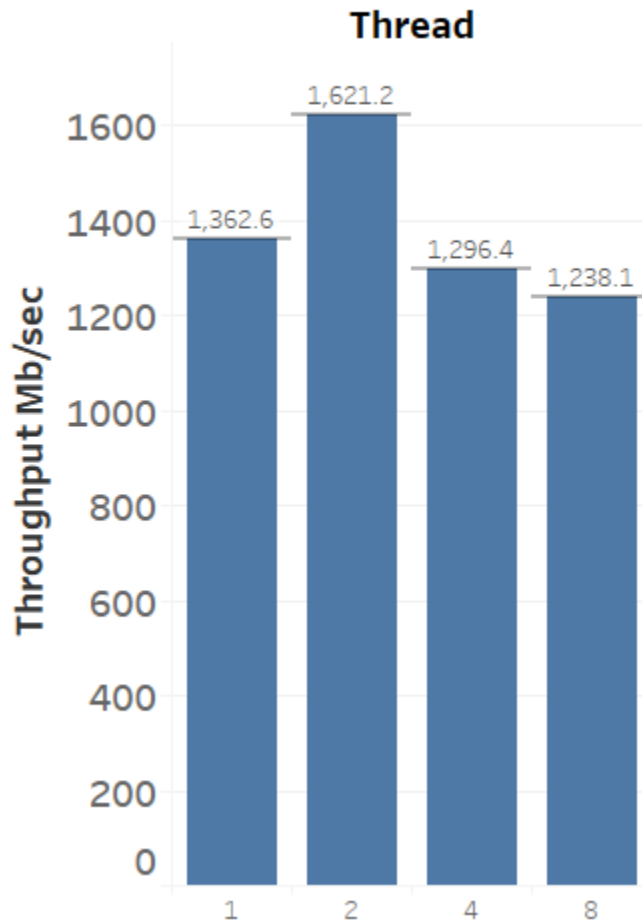
Efficiency is

=3335/4949.4


=67.38%

Network Benchmarking:

TCP



TCP




Above represents throughput of TCP communication over loopback interface card and different number of thread and average and standard deviation of all threads. For this benchmark we fix packet size to transfer(64kb). Throughput is measure in megabits per seconds

We observed that for thread 1 we get low throughput as compare to other and its take more time for transfer the packet from client to server.

Throughput is increases as we increase the concurrency for TCP communication.

TCP




Above represents latency of TCP communication over loopback interface card and different number of thread. For this benchmark we fix packet size to transfer(64kb).


We observed that for thread 1 we get low latency as compare to other and its take more time for transfer the packet from client to server.

We observed that latency decrease as we increase the concurrency for TCP communication.

UDP



UDP




Above represents throughput of UDP communication over loopback interface card and different number of thread and average and standard deviation of all threads. For this benchmark we fix packet size to transfer(64kb). Throughput is measure in megabits per seconds.

We observed that for thread 1 we get low throughput as compare to other and its take more time for transfer the packet from client to server.

We observed that performance for UDP is better than performance of TCP.

UDP



Above represents throughput of UDP communication over loopback interface card and different number of thread. For this benchmark we fix packet size to transfer(64kb).

We observed that for thread 1 we get low throughput as compare to other and its take more time for transfer the packet from client to server.

We observed that latency decrease as we increase the concurrency for TCP communication.

IPRef Output:

TCP

```
Server listening on 5201
-----
Accepted connection from ::1, port 60514
[ 5] local ::1 port 5201 connected to ::1 port 60516
[ ID] Interval          Transfer      Bandwidth
[ 5] 0.00-1.00    sec  2.36 GBytes  20.3 Gbits/sec
[ 5] 1.00-2.00    sec  2.43 GBytes  20.9 Gbits/sec
[ 5] 2.00-3.00    sec  2.34 GBytes  20.1 Gbits/sec
[ 5] 3.00-4.00    sec  2.22 GBytes  19.0 Gbits/sec
[ 5] 4.00-5.00    sec  2.47 GBytes  21.3 Gbits/sec
[ 5] 5.00-6.00    sec  2.26 GBytes  19.4 Gbits/sec
[ 5] 6.00-7.00    sec  2.27 GBytes  19.5 Gbits/sec
[ 5] 7.00-8.00    sec  2.08 GBytes  17.8 Gbits/sec
[ 5] 8.00-9.00    sec  2.07 GBytes  17.8 Gbits/sec
[ 5] 9.00-10.00   sec  2.89 GBytes  24.8 Gbits/sec
[ 5] 10.00-10.04  sec   107 MBytes  25.5 Gbits/sec
-----
[ ID] Interval          Transfer      Bandwidth
[ 5] 0.00-10.04   sec    0.00 Bytes   0.00 bits/sec
[ 5] 0.00-10.04   sec   23.5 GBytes  20.1 Gbits/sec
-----
Server listening on 5201
-----
```

sender
receiver

```
-----
Server listening on 5201
-----
Accepted connection from ::1, port 60514
[ 5] local ::1 port 5201 connected to ::1 port 60516
[ ID] Interval          Transfer      Bandwidth
[ 5] 0.00-1.00    sec  2.36 GBytes  20.3 Gbits/sec
[ 5] 1.00-2.00    sec  2.43 GBytes  20.9 Gbits/sec
[ 5] 2.00-3.00    sec  2.34 GBytes  20.1 Gbits/sec
[ 5] 3.00-4.00    sec  2.22 GBytes  19.0 Gbits/sec
[ 5] 4.00-5.00    sec  2.47 GBytes  21.3 Gbits/sec
[ 5] 5.00-6.00    sec  2.26 GBytes  19.4 Gbits/sec
[ 5] 6.00-7.00    sec  2.27 GBytes  19.5 Gbits/sec
[ 5] 7.00-8.00    sec  2.08 GBytes  17.8 Gbits/sec
[ 5] 8.00-9.00    sec  2.07 GBytes  17.8 Gbits/sec
[ 5] 9.00-10.00   sec  2.89 GBytes  24.8 Gbits/sec
[ 5] 10.00-10.04  sec   107 MBytes  25.5 Gbits/sec
-----
[ ID] Interval          Transfer      Bandwidth
[ 5] 0.00-10.04   sec    0.00 Bytes   0.00 bits/sec
[ 5] 0.00-10.04   sec   23.5 GBytes  20.1 Gbits/sec
-----
Server listening on 5201
-----
```

sender
receiver

UDP

```
Server listening on 5201
-----
Accepted connection from ::1, port 60538
[ 5] local ::1 port 5201 connected to ::1 port 33568
[ ID] Interval          Transfer      Bandwidth      Jitter    Lost/Total Datagrams
[ 5] 0.00-1.00 sec      128 KBytes    1.05 Mbits/sec  0.002 ms   0/6 (0%)
[ 5] 1.00-2.00 sec      128 KBytes    1.05 Mbits/sec  0.011 ms   0/6 (0%)
[ 5] 2.00-3.00 sec      128 KBytes    1.05 Mbits/sec  0.016 ms   0/6 (0%)
[ 5] 3.00-4.00 sec      128 KBytes    1.05 Mbits/sec  0.018 ms   0/6 (0%)
[ 5] 4.00-5.00 sec      128 KBytes    1.05 Mbits/sec  0.023 ms   0/6 (0%)
[ 5] 5.00-6.00 sec      128 KBytes    1.05 Mbits/sec  0.022 ms   0/6 (0%)
[ 5] 6.00-7.00 sec      128 KBytes    1.05 Mbits/sec  0.016 ms   0/6 (0%)
[ 5] 7.00-8.00 sec      128 KBytes    1.05 Mbits/sec  0.018 ms   0/6 (0%)
[ 5] 8.00-9.00 sec      128 KBytes    1.05 Mbits/sec  0.020 ms   0/6 (0%)
[ 5] 9.00-10.00 sec     128 KBytes    1.05 Mbits/sec  0.021 ms   0/6 (0%)
[ 5] 10.00-10.04 sec     0.00 Bytes    0.00 bits/sec   0.021 ms   0/0 (0%)
-----
[ ID] Interval          Transfer      Bandwidth      Jitter    Lost/Total Datagrams
[ 5] 0.00-10.04 sec     0.00 Bytes    0.00 bits/sec   0.021 ms   0/60 (0%)
-----
Server listening on 5201
-----
```

```
Connecting to host localhost, port 5201
[ 4] local ::1 port 60516 connected to ::1 port 5201
[ ID] Interval          Transfer      Bandwidth      Retr      Cwnd
[ 4] 0.00-1.00 sec      2.46 GBytes    21128 Mbits/sec  0         320 KBytes
[ 4] 1.00-2.00 sec      2.43 GBytes    20850 Mbits/sec  0         320 KBytes
[ 4] 2.00-3.00 sec      2.34 GBytes    20076 Mbits/sec  0         320 KBytes
[ 4] 3.00-4.00 sec      2.23 GBytes    19124 Mbits/sec  0         320 KBytes
[ 4] 4.00-5.00 sec      2.47 GBytes    21253 Mbits/sec  0         320 KBytes
[ 4] 5.00-6.00 sec      2.24 GBytes    19257 Mbits/sec  0         320 KBytes
[ 4] 6.00-7.00 sec      2.29 GBytes    19656 Mbits/sec  0         320 KBytes
[ 4] 7.00-8.00 sec      2.08 GBytes    17840 Mbits/sec  0         320 KBytes
[ 4] 8.00-9.00 sec      2.08 GBytes    17855 Mbits/sec  0         320 KBytes
[ 4] 9.00-10.00 sec     2.89 GBytes    24829 Mbits/sec  0         320 KBytes
-----
[ ID] Interval          Transfer      Bandwidth      Retr
[ 4] 0.00-10.00 sec     23.5 GBytes    20187 Mbits/sec  0
[ 4] 0.00-10.00 sec     23.5 GBytes    20187 Mbits/sec
sender
receiver

iperf Done.
```

Bandwidth for IPRef benchmark for TCP is 20.1 Gbits/sec. We got 1621 Mbits/s.

Efficiency for TCP is

$$= 1621 / (20.1 * 1000)$$

$$= 8.06\%$$

Bandwidth for IPRef benchmark for UDP is 23.5 Gbits/sec. We got 1927 Mbits/s.

Efficiency for UDP is

$$= 1927 / (23.5 * 1000)$$

$$= 8.2\%$$